

# Comparative Analysis of RoBERTa and CodeBERT for Automated Misconfiguration Detection in Dockerfiles

(MIT Thesis Viva Voice Examination)



Tribhuvan University  
Institute of Science and Technology  
**Central Department of Computer Science & Information Technology**  
**(CDCSIT TU)**

**Presented By:**

Deepak Aryal

Roll no: 7925005

CDCSIT TU

**Supervised By:**

Asst. Prof. Bikash Balami

CDCSIT TU

Date: January 23, 2026

# ABSTRACT

- This thesis explores automated detection of Dockerfile misconfigurations using transformer-based NLP models.
- A two-stage framework is proposed for binary and rule-specific misconfiguration classification.
- The models are evaluated on a dataset of public Dockerfiles.
- Results show high detection accuracy, with CodeBERT slightly outperforming RoBERTa.

# Outline

- Introduction
- Problem Statement
- Objectives
- Scope and Limitation
- Literature Review
- Methodology
- Dataset Description
- Data Preprocessing
- Description of Algorithms
- Implementation Details
- Result Analysis
- Inference
- Conclusion
- Future Recommendation
- References

# Introduction

- Docker is a core technology in modern DevOps for building and deploying applications
- Dockerfiles define container build processes but are complex and under-analyzed
- Misconfigurations in Dockerfiles can cause security vulnerabilities and build failures
- Existing tools detect rule violations but fail to generalize to complex patterns
- This research evaluates transformer-based NLP models to improve Dockerfile misconfiguration detection

This naturally leads to the problem addressed in this thesis.

# Problem Statement

- Public Dockerfiles frequently violate best practices and contain security risks
- Many engineers lack deep security expertise when writing Dockerfiles
- Existing tools are rule based and limited to predefined checks which fail to detect unseen patterns
- NLP/AI for Dockerfiles is still underexplored

# Problem Statement

## traefik:latest [v2.11.34]

- Traefik http router rule



```
traefik.http.routers.${PROJECT}.rule=Host(`api.${DOMAIN}`) ||  
(HostRegexp(`{subdomain:(app)}.${DOMAIN}`) && PathPrefix(`/api/`))
```

## traefik:latest [v3.6.6]

- Traefik http router rule



```
traefik.http.routers.${PROJECT}.rule=Host(`api.${DOMAIN}`) ||  
(HostRegexp(`{subdomain:(app)}.${DOMAIN}`) && PathPrefix(`/api/`))
```



```
traefik.http.routers.${PROJECT}.rule=Host(`api.${DOMAIN}`) || ((Host(`app.${DOMAIN}`) && PathPrefix(`/api/`))
```

# Objectives

- To design and develop automated detection framework using either RoBERTa or CodeBERT, and find out which one has better performance
- To evaluate and compare the performance of RoBERTa and CodeBERT for automated misconfiguration detection in Dockerfiles, using metrics such as accuracy, precision, recall, F1-score, and AUC-ROC

# Scope and Limitation

 **Models compared:** RoBERTa & CodeBERT

 **Objective:** Automated detection of Dockerfile misconfigurations

 **Analysis type:** Static, NLP-based evaluation

 **Performance metrics:** Accuracy | Precision | Recall | F1-score | AUC-ROC

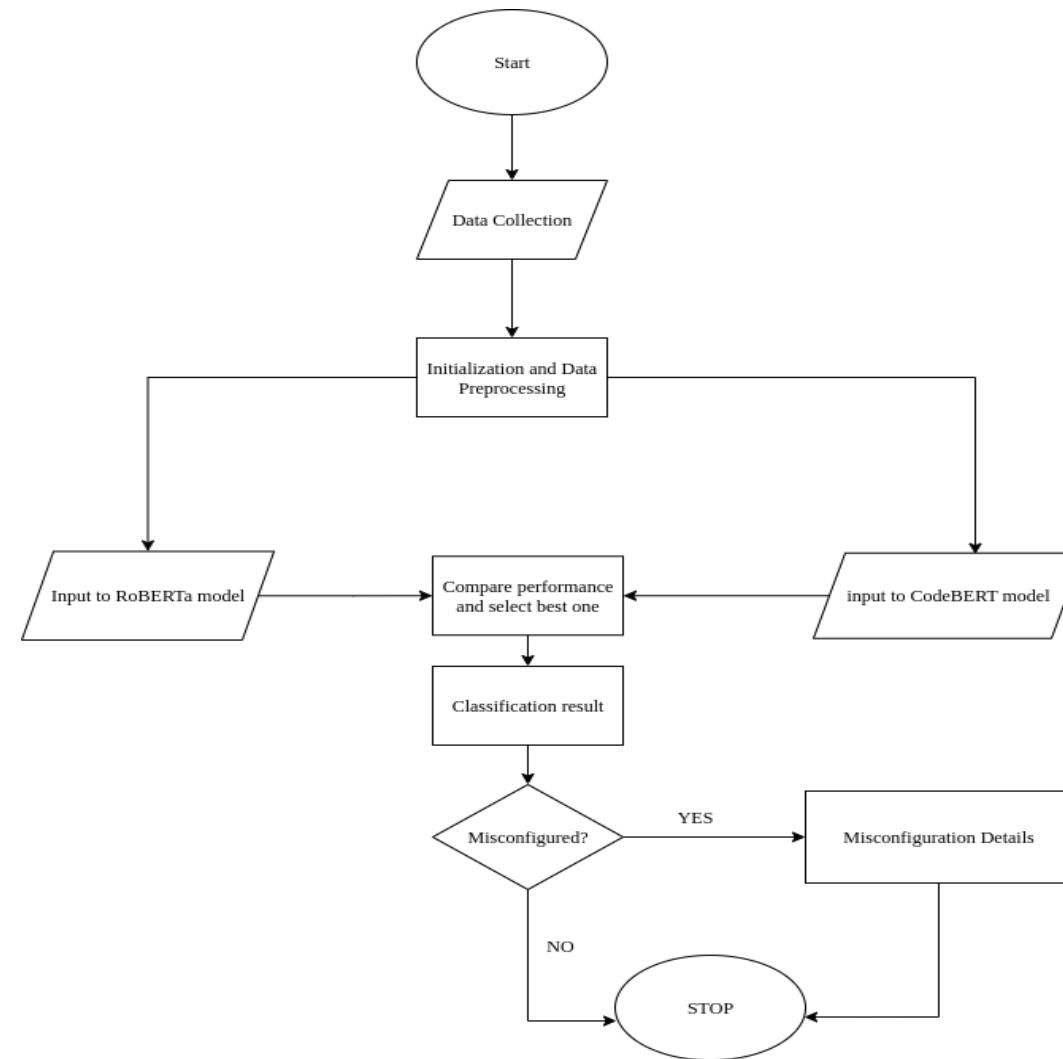
-  **Scope:** Only Dockerfiles
-  **Dataset:** Publicly available Dockerfiles
-  **Models:** Only RoBERTa and CodeBERT
-  **Resources:** Limited computational power may affect experiments and tuning

# Literature Review

Study	Dataset	Key findings	Research gap
Learning from, DevOps Artifacts	Expert & non-expert Dockerfiles	Non-expert Dockerfiles violate best practices $\sim 6\times$ more; IDE integration suggested	Needs adaptive AI methods for complex misconfigs
Shipwright	Faulty Dockerfiles	Automated repair suggestions; Improves build correctness	Focus on build failure, not misconfiguration
Dockerfile Flakiness	Flaky Dockerfiles	Characterizes flakiness; Improves repair success	No line-level security detection
Dockerfile Linting	Docker Workflows	Ensures consistent quality via linting	No AI/NLP misconfiguration detection
Not all Dockerfile Smells are the Same	39,000 Dockerfiles	Experts overlook smells; Validates linting	No AI/NLP detection explored
DRIVE: Rule Mining	Dockerfiles	Identifies semantic & syntactic rules; Improves quality	Focus on static rules, not AI/NLP

# Methodology

## System flow chart



# Dataset Description

- Source: Publicly available Dockerfiles from GitHub
- Primary Corpus: Henkel et al. Dockerfile corpus (ICSE 2020)
  - ~178,000 real-world Dockerfiles
- Extended data: Additional Dockerfiles collected via GitHub Search API
- Processing and labels: Dockerfiles parsed line by line; each instruction (e.g., FROM, RUN, COPY) labeled as correct or wrong for training and evaluation of NLP models.
- Dataset Size:
  - ~178,800 Dockerfiles
  - ~1.6 million Dockerfile instructions

# Data Preprocessing

## Model training strategy

- Stage 1: Binary misconfiguration classification
  - Correct vs wrong
- Stage 2: Rule-specific misconfiguration classification
  - Misconfigured instructions with rule id

# Data Preprocessing

## Data splitting

- Stratified splitting into:
  - Training: 80%
  - Validation: 10%
  - Test: 10%
- Separate strategies for stage 1 and stage 2
- Due to computational resource constraints, a representative subset of the dataset was used for training and evaluation.

# Data Preprocessing

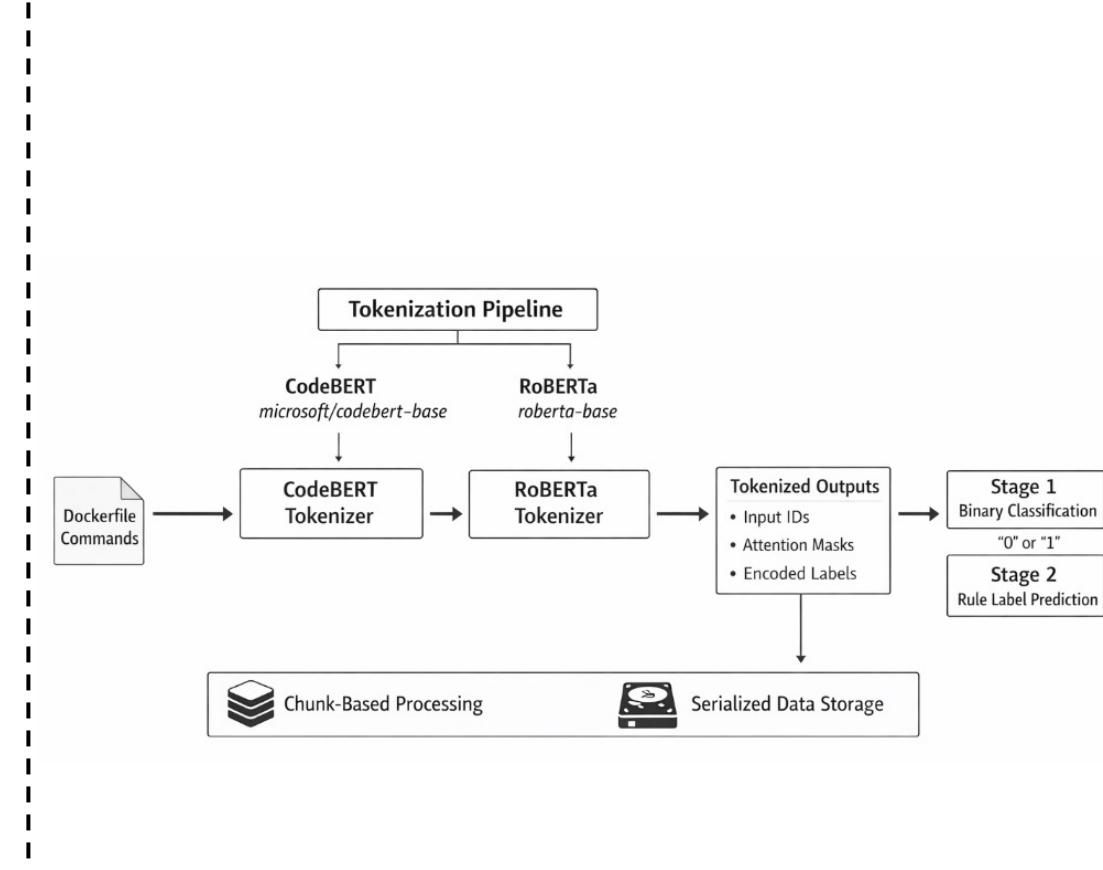
## Data splitting

Aspect	Stage 1: Binary classification	Stage 2: Rule-level classification
Task	Correct vs Wrong	Rule-level misconfiguration classification
Initial data	~1.6M instructions	Misconfigured instructions only
Data reduction and balancing method	Downsampled and selected 200K samples per class	Rare rules removed (<200 samples) and capped dominant rules (<=5K)
Final split	80% / 10% / 10% (stratified)	80% / 10% / 10% (rule-stratified)
Motivation	Computational feasibility	Fair multi-class learning and fair computational feasibility

# Data Preprocessing

## Tokenization and encoding

- Unified subword tokenization applied for both learning stages
- Fixed input length of 256 tokens with truncation and padding
- Training set only label encoding, with pre-tokenized inputs stored for efficient reuse

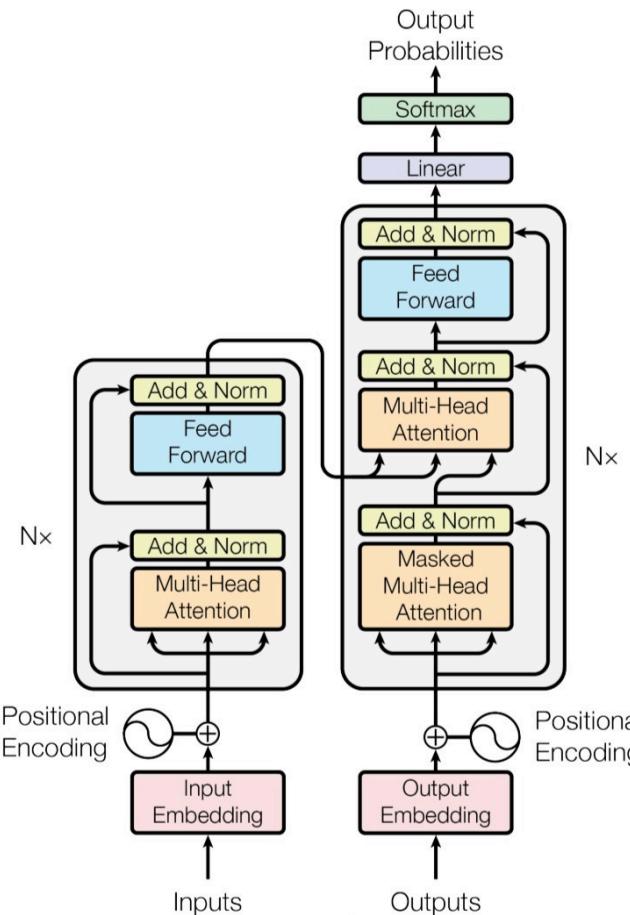


# Data Preprocessing

## Dataset summary

Aspect	Stage 1: Binary classification	Stage 2: Rule-level classification
Initial dataset	~1.6M instructions	Same dataset
Classes	2 (correct vs wrong)	63 rules
Class balance	2,00,000 per class	Filtered 200+ samples, capped at 5,000 per rule
Total samples used	4,00,000	1,73,297
Training set (80%)	3,20,000	1,38,637
Validation set (10%)	40,000	17,330
Test set (10%)	40,000	17,330
Purpose	Detect misconfiguration	Predict violated rule if misconfigured

# Description of Algorithms



- Faster than RNNs/LSTMs: parallel training on large datasets
- Attention mechanism: focuses on most relevant words in context
- Encoder/Decoder: encoder for classification, both for translation
- Self-attention: captures dependencies even across distant words
- Positional encoding: keeps track of word order

# Description of Algorithms

## **RoBERTa and CodeBERT**

- Encoder only transformer architecture
- Optimized version of BERT with larger datasets and better training strategies
- Dynamic masking: different tokens are masked in each epoch

# Description of algorithms

## RoBERTa and CodeBERT

### Transformer encoder

- Stack of self-attention layers producing contextualized token representations.
- Input tokens => embedding vectors => processed by encoder layers.

$$\text{Attention } (Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k})V$$

- Where Q, K, V: query, key, value from embeddings and dk: dimension of key vectors

### Classification head

- The [CLS] token embedding from the encoder is passed through a feed-forward layer to produce a probability distribution over class labels, enabling classification of Dockerfiles as correct or misconfigured.

$$y = \text{softmax}(W \cdot h_{CLS} + b)$$

- Where W, b = trainable parameters and y = probability distribution over classes

# Implementation Details

## Model architecture and input processing

### Pretrained models

- RoBERTa-base (text-oriented)
- CodeBERT-base (code-aware)

### Tokenization

- Model-specific tokenizers
- Subword encoding
- Maximum sequence length: 256 tokens

### Input granularity

- Dockerfile commands processed at line level

### Output heads

- Binary classifier (Stage 1)
- Multi-class classifier (Stage 2)

# Implementation Details

## Training configuration and optimization

Optimizer: AdamW

Learning rate:

- Binary classification:  $2 * 10^{-5}$
- Rule classification:  $5 * 10^{-5}$

Batch size: 32

Epochs: 3

Training strategy:

- Validation-based loss monitoring
- Gradient accumulation

Loss function:

- Binary classification: Cross-entropy loss
- Rule classification: Multi-class cross-entropy loss

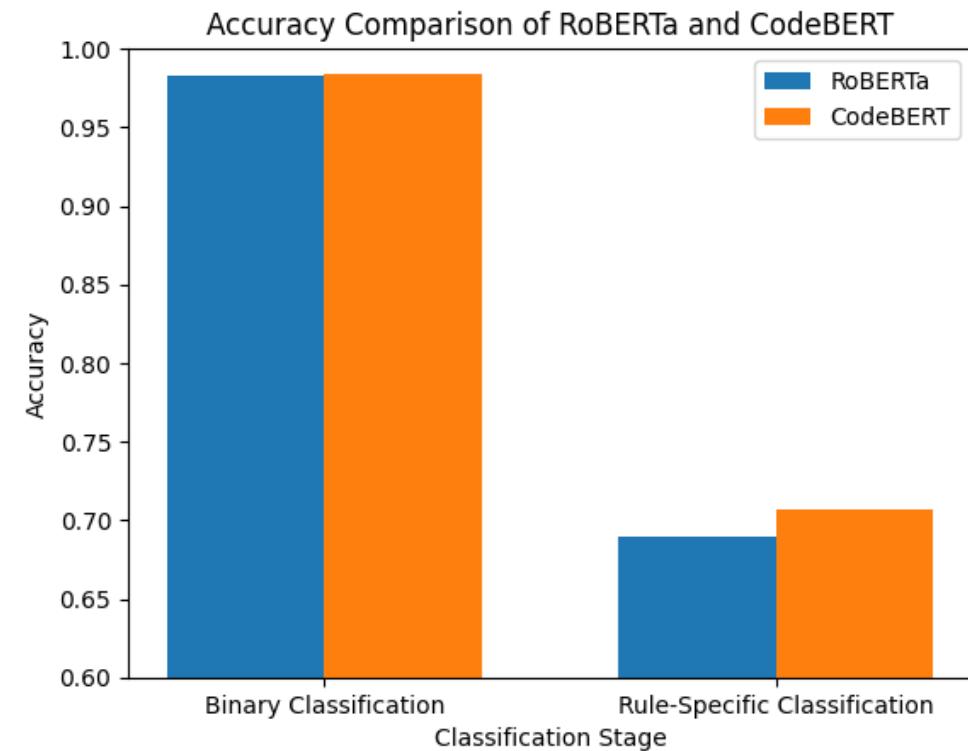
Model persistence:

- Epoch-wise checkpointing for reproducibility

# Result Analysis

## Accuracy

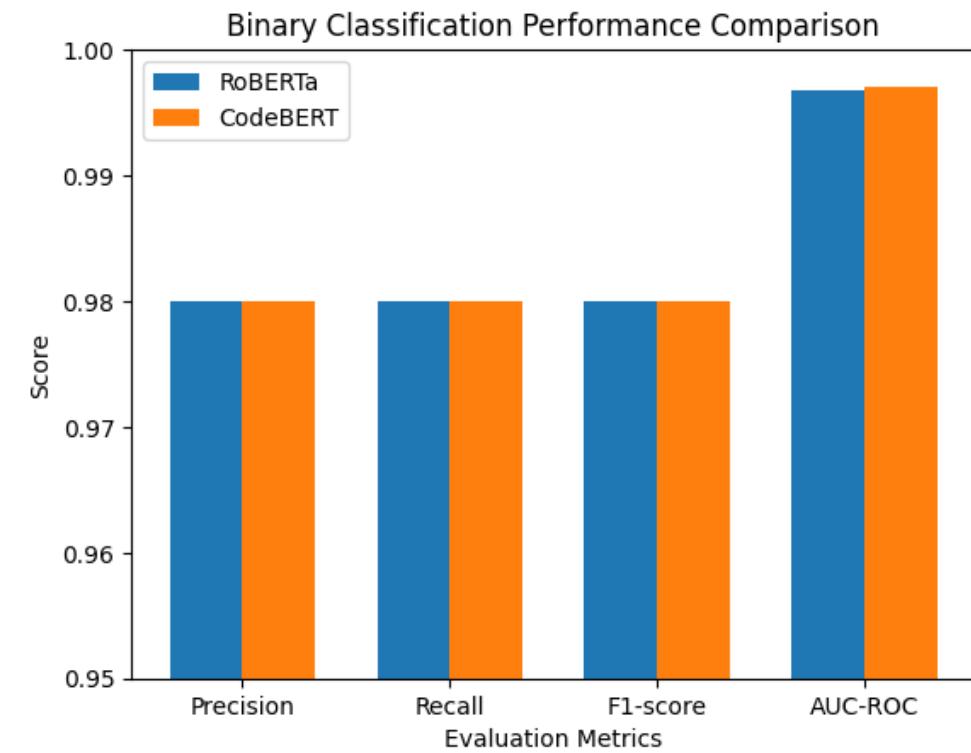
Models \ Stages	RoBERTa	CodeBERT
Binary classification	0.98	0.98
Multi class classification	0.7048	0.7070



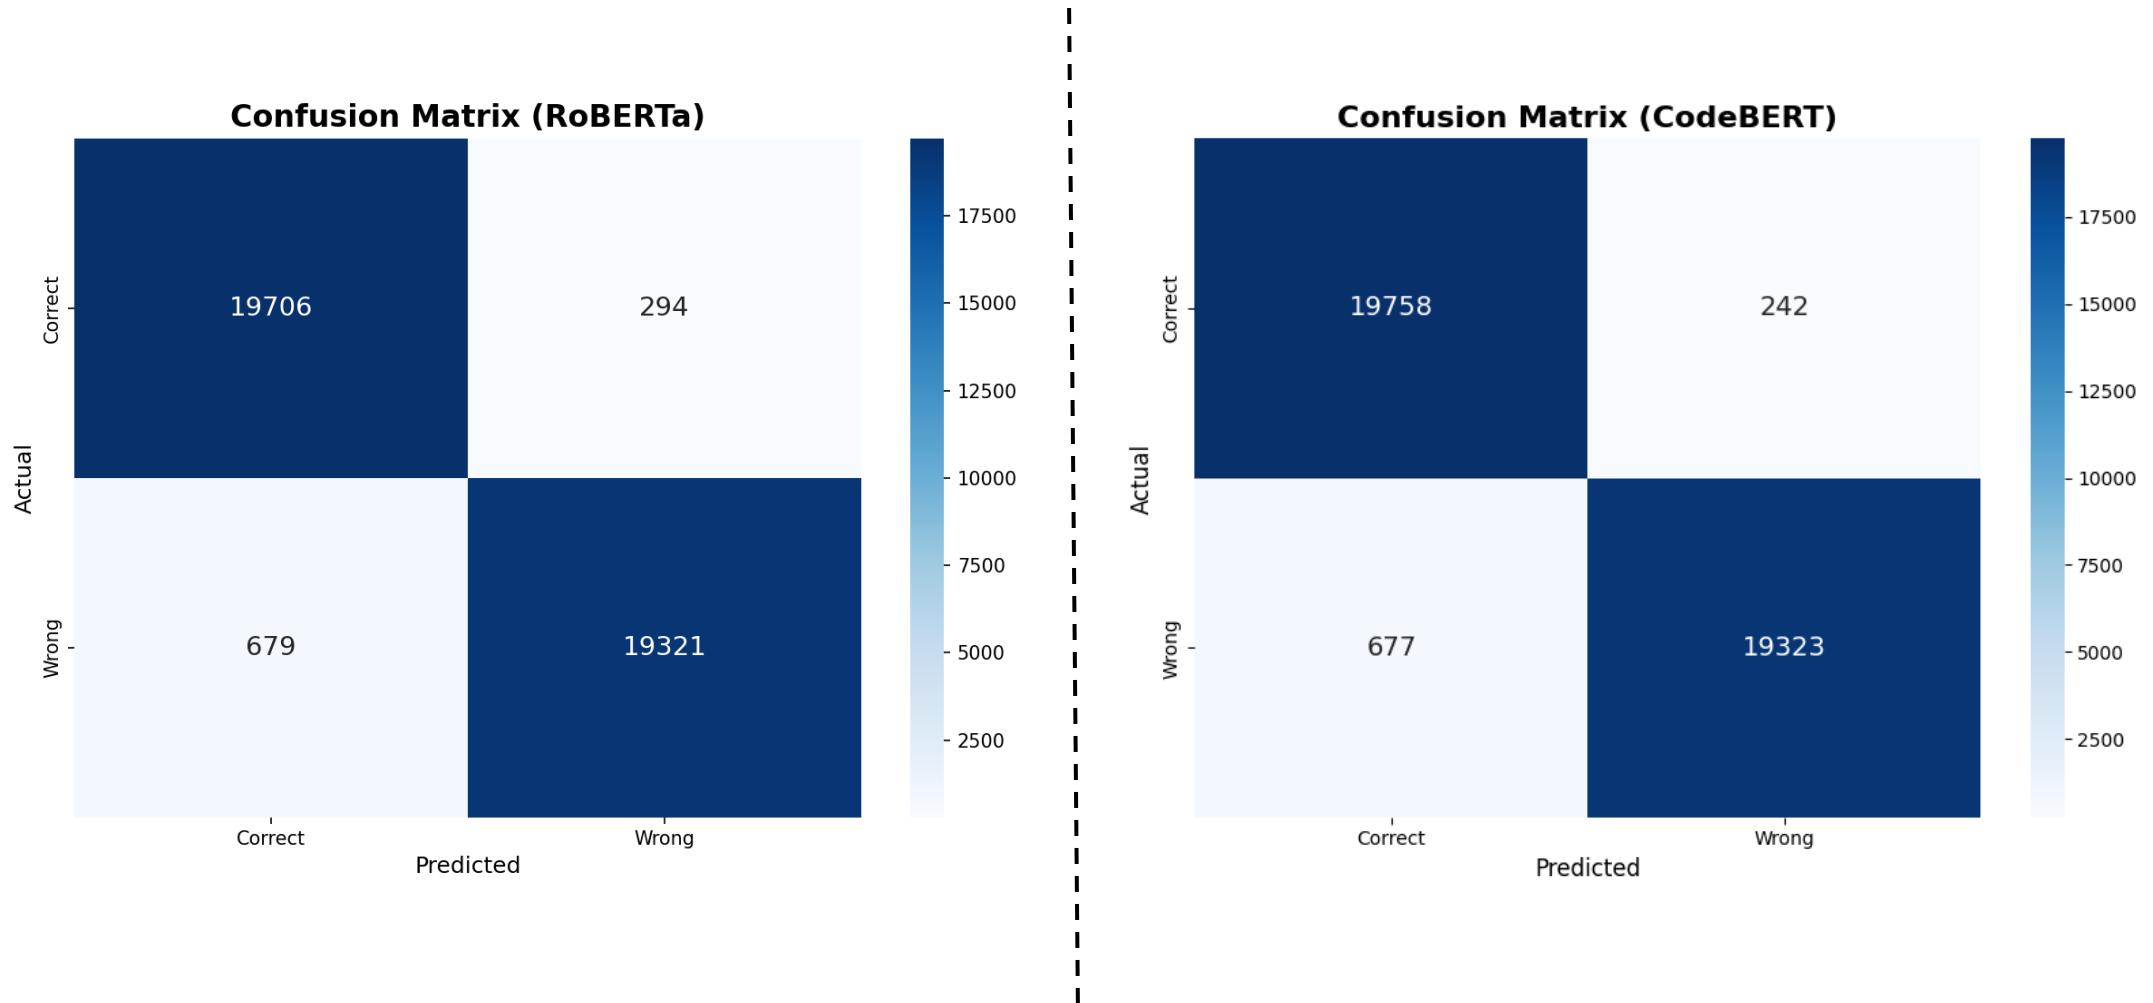
# Result Analysis

## Binary classification performance

Metrics \ Models	RoBERTa	CodeBERT
Precision	0.98	0.98
Recall	0.98	0.98
F1 score	0.98	0.98
AUC-ROC	0.9968	0.9970



# Result Analysis



# Result Analysis

## Rule(Multi) classification performance (Overall)

Metrics	Models	RoBERTa	CodeBERT
Macro-averaged F1		0.7055	0.7136
Weighted-averaged F1		0.7013	0.7040
Top-1 accuracy		0.7048	0.7070
Top-3 accuracy		0.8901	0.8938

# Result Analysis

## Per class rule(multi) classification performance

Metrics	Models	RoBERTa	CodeBERT
Total number of rules		63	63
Average F1-score		0.7055	0.7136
Number of rules with F1 > 0.8		21	21
Number of rules with F1 < 0.5		11	11

# Inference

## RoBERTa

```
=====
RUNNING INFERENCE EXAMPLES ROBERTA
=====

=====
DOCKERFILE LINE ANALYSIS
=====

Input: EXPOSE 8080
-----

[STAGE 1] Checking if configuration is correct...
Prediction: ✓ CORRECT
Confidence: 0.9996 (99.96%)
Inference time: 171.94 ms

[STAGE 2] Skipped (configuration is correct)
-----
```

```
=====
RUNNING INFERENCE EXAMPLES ROBERTA
=====

=====
DOCKERFILE LINE ANALYSIS
=====

Input: RUN apt-get update && apt-get install -y python3
-----

[STAGE 1] Checking if configuration is correct...
Prediction: ✗ WRONG
Confidence: 0.9996 (99.96%)
Inference time: 208.00 ms

[STAGE 2] Identifying violated rule...
Violated Rule: DL3009
Confidence: 0.4662 (46.62%)
Inference time: 1297.62 ms

Total inference time: 1505.62 ms
-----
```

# Inference

## CodeBERT

```
=====
RUNNING INFERENCE EXAMPLES CODEBERT
=====

=====
DOCKERFILE LINE ANALYSIS
=====

Input: EXPOSE 8080
-----

[STAGE 1] Checking if configuration is correct...
Prediction: ✓ CORRECT
Confidence: 0.9998 (99.98%)
Inference time: 163.87 ms

[STAGE 2] Skipped (configuration is correct)
-----
```

```
=====
RUNNING INFERENCE EXAMPLES CODEBERT
=====

=====
DOCKERFILE LINE ANALYSIS
=====

Input: RUN apt-get update && apt-get install -y python3
-----

[STAGE 1] Checking if configuration is correct...
Prediction: ✗ WRONG
Confidence: 0.9996 (99.96%)
Inference time: 164.99 ms

[STAGE 2] Identifying violated rule...
Violated Rule: DL3009
Confidence: 0.4786 (47.86%)
Inference time: 177.43 ms

Total inference time: 342.41 ms
-----
```

# Conclusion

- For binary misconfiguration detection, both CodeBERT and RoBERTa achieve excellent and identical performance ( $\approx 98\%$  accuracy), demonstrating strong and reliable classification capability.
- In rule-specific misconfiguration classification, CodeBERT performs slightly better than RoBERTa, indicating its stronger ability to capture Dockerfile-specific and semantic patterns.
- Model selection depends on application needs: both models are suitable for large-scale binary detection, while CodeBERT is more effective for detailed rule-level misconfiguration analysis.

# Future Recommendation

- Improve rule-specific classification using larger and more diverse datasets, including data augmentation for rare rules
- Explore ensemble models or hybrid approaches combining transformers with rule-based heuristics
- Integrate the proposed models into CI/CD pipelines for real-time misconfiguration detection
- Extend the approach to other IaC technologies such as Terraform, CloudFormation, and Ansible

# References

- [1] S. Phillips, T. Zimmermann, and C. Bird, “Understanding and Improving Software Build Teams,” *Proc. - Int. Conf. Softw. Eng.*, May 2014, doi: 10.1145/2568225.2568274.
- [2] A. Vaswani et al., “Attention is All you Need,” in Neural Information Processing Systems, 2017. [Online]. Available: <https://api.semanticscholar.org/CorpusID:13756489>
- [3] J. Henkel, C. Bird, S. K. Lahiri, and T. Reps, “Learning from, Understanding, and Supporting DevOps Artifacts for Docker,” *2020 IEEE/ACM 42nd Int. Conf. Softw. Eng.*, pp. 38–49, 2020, [Online]. Available: <https://api.semanticscholar.org/CorpusID:211069127>

# References

- [4] J. Henkel, D. Silva, L. Teixeira, M. d'Amorim, and T. Reps, “Shipwright: A Human-in- the-Loop System for Dockerfile Repair,” *2021 IEEE/ACM 43rd Int. Conf. Softw. Eng.*, pp. 1148–1160, 2021, [Online]. Available: <https://api.semanticscholar.org/CorpusID:232045443>
- [5] T. Shabani, N. Nashid, P. Alian, and A. Mesbah, “Dockerfile Flakiness: Characterization and Repair,” *2025 IEEE/ACM 47th Int. Conf. Softw. Eng.*, pp. 1793–1805, 2024, [Online]. Available: <https://api.semanticscholar.org/CorpusID:271855816>
- [6] D. W. Elliott, “Dockerfile Linting: Every time, Everywhere,” 2020, [Online]. Available: <https://medium.com/@david.w.elliott/dockerfile-linting-every-time-everywhere-d8d271a1e650>

# References

- [7] G. Rosa, S. Scalabrino, G. Robles, and R. Oliveto, *Not all Dockerfile Smells are the Same: An Empirical Evaluation of Hadolint Writing Practices by Experts*. 2024. doi:10.1145/3643991.3644905.
- [8] Y. Zhou, W. Zhan, Z. Li, T. Han, T. Chen, and H. C. Gall, “DRIVE: Dockerfile Rule Mining and Violation Detection,” *ACM Trans. Softw. Eng. Methodol.*, vol. 33, pp. 1–23, 2022, [Online]. Available: <https://api.semanticscholar.org/CorpusID:254564701>
- [9] J. Henkel, C. Bird, S. K. Lahiri, and T. Reps, “ICSE 2020 Artifact for: Learning from, Understanding, and Supporting DevOps Artifacts for Docker.” Zenodo, 2020. doi: 10.5281/zenodo.3628771.

**Thank you!!!**